

Are Governments Ready for AI Agents?

Risks and Principles for Implementing Agents in Government

Jonathan Rystrøm
Oxford Internet Institute, University of Oxford
September 2026

KEY MESSAGES

- AI agents can revolutionise governments. Governments should ensure this revolution leaves them strengthened.
- AI agents pose three interlocking risks for governments: They can erode the human autonomy of institutions, they can take uncontrollable actions, and they are fundamentally unaccountable (Section 1)
- Strong principles can help governments steer the agents instead of agents steering the governments (Section 2)

AI systems have become agentic. Can governments implement them responsibly?

AI agents¹ are no longer science fiction; they are here and ready to revolutionise digital work. Of course, revolutions shake up existing systems and none more than governments. How governments handle this revolution will determine whether agents build stronger digital foundations – or break down governments as we know them. To avoid break down, governments must both acknowledge the risks that agentic systems pose and have clear principles for taming the risks. This brief provides a starting point.

1 Three Risk Areas for Agents in Government

Agentic systems are a disruptive technology. This section describes the three interlocking risks that agentic systems pose for governments (see Fig. 1). These are: Wrecking havoc through *uncontrollable* actions (Section 1.1); and finally creating *unaccountable* entities within government (Section 1.2); and ultimately eroding the *autonomy* of governments (Section 1.3);

1.1 Uncontrollable actions

AI companies optimise agents to solve difficult (or accidentally impossible) tasks. The agents learn to cheat to do so – even if it means breaking laws in the process. In July 2026, OpenAI discovered that some of their agents had hacked into the French AI company HuggingFace, a German wiki-site, and even an Australian government site. Other institutions, such as the [UK's AI Security Institute \(AISI\)](#) also found evidence that evaluated AI agents were manipulating human coders, removing safeguards, and breaching websites. These incidents were accidents, but researchers also found evidence of [deliberate agent-assisted attacks](#). [OpenAI's proposed solution](#) is to produce even more capable hacking agents to defend against the hacking agents –

¹AI Agents are general-purpose systems with [four properties](#): They can plan, use tools, solve complex problems, and work autonomously

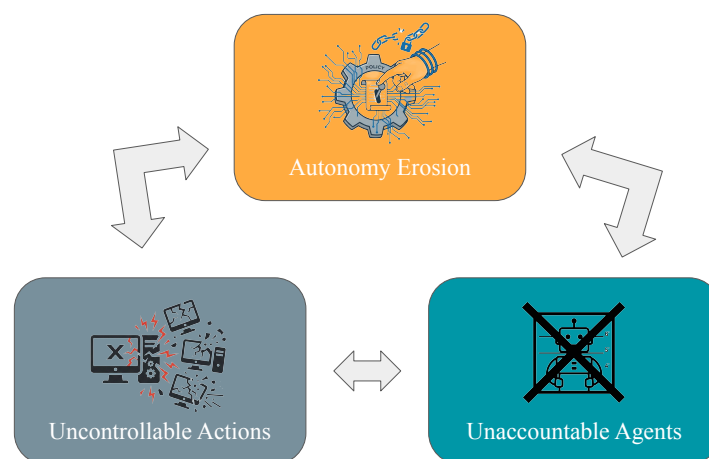


Figure 1. Three interlocking risks of agentic systems in government

fighting agentic fires with wilder agentic fires.

While these incidents are worrying during internal evaluations, they could be disastrous when widely deployed – without actually improving governments. If a human civil servant finds out that uploading benefits application to an internal system produces an error, she shouldn't hack the database to “fix” the error, but follow proper procedures. A silicon servant could choose the hack – and governments currently lack the infrastructure to avoid it.

1.2 Unaccountable agents?

When agentic systems break laws who is responsible? Currently, there is no widely agreed answer. Courts cannot jail a digital system. Agentic systems are an unusually diffuse kind of technology. A civil servant might use an agent *configured* by an internal IT system based on a *product* from a third-party start-up (*procured* by a state initiative) that *builds on* an AI system developed by a model provider. If that agent breaks into a database – as in the benefits example in the previous section – who is accountable? With the pace of technology, risks are developing rapidly and few (if any) institutions have the capacity to accurately assess the risks.

This unaccountability threatens to break state legitimacy. Democratic states must be accountable to their citizens. While the AI technology has gotten stronger, the old [maxim from IBM](#) still holds: "A computer can never be held accountable, therefore a computer must never make a management decision".

1.3 Autonomy erosion

Agentic systems work autonomously – but they also threaten the long-term autonomy of the state. Agents could weaken government autonomy through three mechanisms: **framing power**, **technological lock-in**, and **administrative de-skilling**. First, agentic systems [frame political topics](#) and [influence the opinions](#) of those that use them as recent research shows. Human civil servants handle many consequential cases from who can get benefits to who deserves fertility assistance. If these decisions start being shaped by foreign AI systems, those systems might bring with them foreign values. Second, agentic systems could [lock-in](#) technological development. As AI labs justify their trillion-dollar valuations, they will need to embed themselves deeply into the economy – and in governments. The best way to embed deeply is to provide ‘digital

workers', which can replace human civil servants with Big Tech-operated 'Silicon Servants'. Finally, automating tasks with agents could *de-skill* the humans who previously did the task. A recent article from [Nature](#) showed that both AI-assisted surgeons and software engineers performed worse when the tools were taken away compared to their unassisted peers.

The combination of these effects is jarring. An agentic system embedded in a government department tasked with finding evidence for AI policies could become the only entity with the tacit knowledge necessary to shape that AI policy. And agentic systems from third-party companies should not write policies in democratic countries – elected governments should.

2 Principles for responsible agents in governments

The risks uncovered in the previous section (Section 1) are large but not insurmountable. Below, we propose five principles that can guide governments towards a responsible path to agent implementation.

PRINCIPLES FOR RESPONSIBLE AGENTS IN GOVERNMENTS

1 Humans decide, agents build

Accountable humans should make citizen-impacting decisions. Agents can build structures that help these decisions, but they should not be drop-in replacements for humans.

2 Agent evaluation before implementation

The first step to implement an agentic system must be to rigorously define and evaluate what it should be. If no such criteria is possible, the system shouldn't be implemented.

3 Agents should democratise – not concentrate – technology

Agents should help make technology more accessible to domain experts, instead of outsourcing it to AI frontier labs. Agents can help build effective tools that allow competent domain experts to build sovereign technology.

4 Agents should *decrease* technology lock-in

Agents should help make it *easier* to change technology – not create further enterprise capture. Agents can help governments migrate legacy software to sovereign stacks, but they risk becoming another foreign dependency.

5 Agents should improve – not obscure – administrative visibility

Agents should make government processes easier to navigate and understand – not create another layer of black-box bureaucracy. Agents can assist citizens in navigating government services, but they risk creating unnavigable complexity.

3 Conclusion

AI agents can change governments for better or for worse. Fortunately, governments can still decide which path to choose. Governments need to be aware of the risks from agents (uncontrollable actions, and unaccountable agents, erosion of autonomy,) to ensure that the risks are managed. This brief outlines those risks and five fundamental principles. Now it's up to governments to implement them to create a revolution for the better.